name
email | phone
linkedin | github| portfolio | San Francisco, CA

EDUCATION
San Francisco State University   San Francisco, California
B.S. in Computer Science
- GPA: 3.96/4.00, *Dean's List*   Expected Graduation 2027

EXPERIENCE
AI Engineer Intern, Company, San Francisco, CA   *May 2025 - Aug 2025*
- Contributed with the team to finetune a Qwen-2.5-7B-param using Huggingface, PyTorch, Lambda (Cloud computing), LoRA (cost and memory efficient training) on Q&A tasks and hosted on Cloud Run (Google Cloud Platform).
- Led Dataset curation using pandas, numpy and datasets library
- Evaluated model and hyperparameters tuned to achieve very low valuation loss, tracked using wandb (model logging and experiment tracking).
- Additionally, Developed an AI-powered Outreach Agent using Langchain, Exa.ai along with OpenAI API Integration to automate messaging workflows.

ML Engineering Intern, Compony, Remote, TX   *Oct 2025 – Present*
- Working with the team to train and manage AI Models using Python, PyTorch and transformers; along with model tracking with wandb as well as model serving using endpoints on Azure ML Studio.
- Created a fully working Web Crawler using Scrapy and a scraper using BeautifulSoup4, providing company relevant data from open websites under proper ethics.

PROJECTS
Qwen-2.5-0.5B Finetune Github | Huggingface | Dockerhub
- Independently fine-tuned Qwen-2.5-0.5B using Hugging Face Transformers, PyTorch, LoRA, and DPO (post-training human alignment) on Google Colab A100 (GPU compute) for instruction-following tasks.
- Trained with bf16 (*lower memory usage*), gradient checkpointing, Flash Attention (*faster training*), and tf32 *(for memory efficiency and faster inference)*; experiments tracked in Weights & Biases *(experiment logging)*.
- Packaged an inference-ready Docker image powered by vLLM *(faster inference)*; artifacts published on DockerHub and mirrored on Hugging Face Hub *(deployment-ready)*

GatorGPT Github | Huggingface
- Engineered a 63M parameter transformer model using PyTorch and modern architecture components such as GQA, RoPE, and SwiGLU MLP layers, trained on the TinyStories dataset.
- Deployed and served using vLLM, with the complete model available on Hugging Face for one-click usage.
- Planned next phase involves fine-tuning on university-specific datasets using Direct Preference Optimization (DPO) and Reinforcement Learning *(for personalized alignment after supervised fine-tuning)*.

Max – Personal Voice Assistant Github
- Developed a voice activated AI Agent using Langchain, OpenAI, and SpeechRecognition to automate tasks along with hands-free interaction.
- Tools like, Web Search, Youtube Streaming, Emailing, File generation with wide range of extensions, knowledge based Retrieval Augmented Generation, controlling camera and much more are integrated.

And more on Github.